

I trained a neural network on 214,425 bank-quarter examples.

Among the original four forecasts, predicting zero growth had the lowest mean absolute error.

PREDICT 0%


for every bank.

Zero: 3.601 pp MAE | MLP: 3.627 | Ridge: 3.641

MAE measures the average absolute miss in percentage points. All scores use the same 13,532 bank-quarter examples in three already-examined 2024 quarters. Follow-up baselines appear later.

Many training targets sit near zero.



That makes a constant forecast a useful benchmark. Its later performance still has to be measured.



This histogram shows the central 98% of training log-growth targets. The full tails remain in the experiment. The horizontal unit is $100 \times \log$ growth; it is not ordinary percentage-point growth.

The error measure changes the winner.


Original four models; lower is better. Zero leads MAE. The neural network leads RMSE.



MAE differences among zero, Ridge and MLP are only about 0.04 pp. RMSE gives large misses more weight. These rankings describe the historical sample; they do not establish future superiority.

Zero growth gives every bank the same forecast.

Imagine ten banks and enough time to examine only one. The zero-growth forecast leaves all ten tied.




**ONE REVIEW SLOT.
TEN IDENTICAL SCORES.**

Zero growth remains a valid point-forecast baseline. Selecting among tied forecasts requires a separate rule. Random selection supplies the comparison for the review-list experiment.

Ridge selected 112 banks that later had the lowest deposit growth.


Imagine you can examine 454 of 4,536 banks.
You want to find banks whose deposits will grow least—or shrink most—next quarter.
Ridge’s list included 112 of that quarter’s lowest-growth 454 banks; chance would find about 45.



March → June 2024. Review capacity is hypothetical. Lowest growth can still be positive. Dots show rounded shares; each dot does not represent one bank.


Deposit size alone found about 19 cases per 100 selections.

A case means a selected bank that later fell in its quarter’s lowest-growth 10%.
Using all five inputs raised the count to about 22 per 100 selections.
Random selection would find about 10. Size alone is a serious benchmark.



The calendar earns a place in the next frozen test.


Training-period decline rates differ by predictor quarter. A seasonal-median forecast uses only training data.



Seasonal median: 3.500 pp MAE; zero growth: 3.601 pp on reused 2024 data. This baseline was added after 2024 had been examined. Its fitting uses earlier rows; its development was post-hoc.

These pretrained models did not beat the simple benchmarks.

Same evaluation rows. The small neural network had lower RMSE than either foundation model in this experiment.



As-is means frozen pretrained weights. Adaptation updated only the forecast head for 256 sampled steps, with earlier validation selecting the checkpoint. Pretraining overlap is unknown. This limited test does not establish what other checkpoints or training budgets could achieve.