

I trained a neural network on 214,425 bank-quarter examples.

And the lowest average miss among the original four forecasts came from this:

PREDICT 0%


for every bank.

Zero: 3.601 pp MAE | MLP: 3.627 | Ridge: 3.641

MAE measures the average absolute miss in percentage points. All scores use the same 13,532 bank-quarter examples in three already-examined 2024 quarters. Follow-up baselines appear later.

Wait. How can predicting zero win anything?



Look where the training examples pile up. Zero is worth testing. The later scores tell us whether it works.



This histogram shows the central 98% of training log-growth targets. The full tails remain in the experiment. The horizontal unit is $100 \times \log$ growth; it is not ordinary percentage-point growth.

Okay, so which forecast actually wins?


Average size of the miss: zero. Give the biggest misses more weight: the neural network. Same banks.



MAE differences among zero, Ridge and MLP are only about 0.04 pp. RMSE gives large misses more weight. These rankings describe the historical sample; they do not establish future superiority.

Good. Now tell me which bank to examine first.

Bank A gets zero. Bank B gets zero. Every bank gets zero. Imagine you have time to examine just one.






**So... who goes first?
Everybody is tied.**

Zero growth remains a valid point-forecast baseline. Selecting among tied forecasts requires a separate rule. Random selection supplies the comparison for the review-list experiment.

Random gets about 45. Ridge got 112.

Imagine 4,536 bank reports on your desk. You have time to examine 454.
So you want a list that finds banks whose deposits will grow least, or shrink most.
Three months later, count how many selected banks landed in the lowest-growth 454.




March → June 2024. Review capacity is hypothetical. Lowest growth can still be positive. Each full grid represents 454 selections.
Colored dots show rounded shares, not individual banks.

Okay, but how much of this is just bank size?

Give Ridge just one clue: deposit size. It finds about 19 per 100 selections.

Give it all five inputs: about 22. That's surprisingly close.


Here, 'finds' means the selected bank later landed in its quarter's lowest-growth 10%.



Equal-quarter averages across three reused 2024 quarters; select the lowest predicted-growth 10%. Size-only Ridge was added after examining 2024. The observed gap needs a fresh test; no operational savings were measured.

So what if the forecast knows what quarter it is?

Training years: declines were more common after Q1. So predict the training median growth for each quarter of the year.




And... the seasonal rule beat zero on MAE: 3.500 vs 3.601 pp.

Seasonal median: 3.500 pp MAE; zero growth: 3.601 pp on reused 2024 data. This baseline was added after 2024 had been examined. Its fitting uses earlier rows; its development was post-hoc.

Surely a pretrained model can beat zero?

On RMSE, these tested versions did not. Even after updating their forecasting heads. Lower is better.



Next: freeze the rules. Save new forecasts. Then wait for the answers.

As-is means frozen pretrained weights. Adaptation updated only the forecast head for 256 sampled steps, with earlier validation selecting the checkpoint. Pretraining overlap is unknown. This limited test does not establish what other checkpoints or training budgets could achieve.