




NEURAL NETS LEARN COMPLEX FUNCTIONS. MINE LOST MAE TO A CONSTANT.

Predicting next-quarter deposit growth from quarterly bank reports.





Actual predictions. Both axes: -10.0% to 22.4% (actual-growth central 98%); 284 rows outside either axis omitted in both panels.
All 13,532 rows scored. 214,425 training examples. Original four-model comparison; historical results.

WHY WAS ZERO SO HARD TO BEAT?



Central 98% of training log-growth targets. Full tails stayed in the experiment.
This concentration motivates the benchmark; it does not by itself explain the later score.

ZERO WON MAE. THE NET WON RMSE.




Different scoring metric. Different winner.

Original four models; same evaluation rows. Errors are in ordinary-growth percentage points. Lower is better. Small historical differences do not establish future superiority.


ZERO CAN FORECAST. IT CAN'T RANK.

Every bank → 0% predicted growth → complete tie



March → June 2024. 4,536 eligible banks; 454 hypothetical review slots.
Found = later in that quarter's lowest-growth 10%, which can include positive growth.

**I THREW OUT 4 INPUTS.
KEPT ONLY DEPOSIT SIZE.**





Same model family. Four features removed. 3.1 percentage-point difference.

Equal-quarter average across three reused 2024 quarters. Size-only challenger added after examining 2024. This comparison does not decompose the five-feature model's causes.

A CALENDAR RULE BEAT ZERO ON HISTORICAL MAE.

How often did deposits fall next quarter?



OKAY. THAT GETS A FRESH TEST.

Seasonal median: predict each calendar quarter’s training-period median growth. Developed after examining 2024. Historical comparison, not fresh confirmation. Analysis notebook built with wm-notecards.